



# Future Directions in DSL Research

---

**Hassan Chafi**, Arvind Sujeeth, Kevin Brown,  
Sungpack Hong, Zach Devito and Kunle Olukotun  
**Oracle Labs**  
**Stanford University**

# Required Support For Multiple DSLs

---

- Each new DSL supported by Delite introduces some new requirement
  - New Ops
  - New optimization requirements
- The good news is that we are finding a lot of re-use between different DSLs which validates the need for an infrastructure

# DSLs Implemented So Far

---

- OptiML
  - Original DSL
- OptiQL
  - LINQ like on Delite
- OptiGraph
  - Green-Marl on Delite
- OptiMesh
  - Liszt on Delite
- OptiCollections
  - Scala Collections on Delite

# Re-Use Across DSLs

---

| DSL             | Delite Ops                                | Generic Optimizations         |
|-----------------|---|-------------------------------|
| OptiML          | Map, ZipWith, Reduce, Foreach, Hash, Sort | CSE, DCE, code motion, fusion |
| OptiQL          | Map, Reduce, Filter, Sort, Hash, Join     | CSE, DCE, code motion, fusion |
| OptiGraph       | ForEachReduce, Map, Reduce, Filter        | CSE, DCE, code motion         |
| OptiMesh        | ForEachReduce                             | CSE, DCE, code motion         |
| OptiCollections | Map, Reduce, Filter, Sort, Hash, ZipWith  | CSE, DCE, code motion, fusion |

# Restrictions & Semantics

---

- The domain-specific restrictions and semantics are currently implementing in an adhoc fashion by DSL author
- Some restrictions and semantics of DSL operations are currently provided by the infrastructure
  - Side-effect management
- Can we generalize and provide a framework for DSL authors to specific the semantics and restrictions imposed by a DSL declaratively?

# DSL Extensibility

---

- Need to be able to start with a DSL or part of a DSL and extend it with some extra domain-specific optimizations and constructs
- How does extensibility interact with the restrictions and semantics of the DSLs?

# OptiML extends OptiLA

---

- **trait OptiLA extends**  
**OptiLAScalaOpsPkg with LanguageOps**  
**with ArithOps with VectorOps //**  
**with ...**
- **trait OptiML extends OptiLA with**  
**OptiMLScalaOpsPkg with**  
**OptiMLVectorOps with StreamOps //**  
**with ...**

# DSL Interoperability

---

- How does the DSL interoperate with the host language?
  - Currently we lift the host language into the DSL
  - The non-lifted parts are staged away during the first part of compilation
- Interesting applications will also use multiple DSLs
- How should these DSLs interoperate?
- At what level of Granularity?



# OptiML and OptiQL interoperability

---

```
val orderData = DeliteRef[Array[Record]]()
val theta = DeliteRef[Array[Double]]()
OptiQL {
  // customers: Array[Customer], orders: Array[Order]
  val orders = customers Join(orders)
  WhereEq(_.c_custkey, _.o_custkey)
  Select((c,o) => new Result {
    val nationKey = c.c_nationkey
    val acctBalance = c.c_acctbal
    val price = o.o_totalprice
  })
  orderData.set(orders)
}
OptiML {
  // run linear regression on price
  val data = Matrix(orderData.get.map(t =>
    Array(t.getDouble(1),t.getDouble(2),t.getDouble(3))))
  val x = data.sliceCols(0,1)
  val y = data.getCol(2)
  theta.set(linreg.weighted(x,y).toArray)
}
println("theta: " + theta.get)
```

# Abstracting Data Structures

---

- Seen how to abstract code in Delite and retarget it to different hw resources
- Need to do the same for data representation
  - current status of stable Delite version is that the DSL author is responsible for providing concrete implementations of the back-end data structures

# Data Structures in Delite

---

- **Current:** DSL developers implement their own data structures for each target device
- **New:** everything is a Struct/Record, specified programmatically
  - instantiation and field access lifted into IR

```
def Complex(re: Rep[Double], im: Rep[Double]) =  
  new Record { val real = re; val imag = im }  
  //lifted to __new and forwarded to Delite  
  
val x = Complex(0,0)  
x.real //type-checked that field "real" exists, then forwarded to Delite
```

# Why Records?

---

- We can auto-generate the back-end implementation to different platforms
  - Supported by large number of target languages
- We can reason about which part of the record is actually used
  - Unwrap the record and just pass around required fields in generated code
  - Unused fields can be eliminated all together
- We can perform automatic AoS -> SoA conversion
  - Instantiate only arrays of primitives in the generated code

# AoS -> SoA Optimization

---

- Provide familiar AoS form to the DSL developer
- Perform SoA transformations transparently for DeliteOps
  - Functions returning record types split into result for each component
  - Create separate loop to compute each component
  - Unused components are dead-code-eliminated
  - Loop fusion recombines live components into single loop

# Integrating with Control Flow

```
def conj(c: Rep[Complex]) = Complex(c.real, -c.imag)

//a: Rep[Array[Complex]]
val b = if (x > 7) a.map(conj) else a
```

- 'a' exists in IR as ("real" -> Rep[Array[Double]], "imag" -> Rep[Array[Double]])
- a.map is split into map for each component
  - map over real component is optimized away (identity function)
- real array is unaffected by conditional, simply re-used in generated code

# Abstracting DSL Analysis and Transformation

---

- Delite currently supports abstractions for parallel execution patterns
- Delite supports optimizations on IR nodes it understands
  - This benefits DSL nodes as well
- Support for Domain-Specific analysis and IR transform in Delite is limited to IR creation time
- More complex use cases are handled by the DSL author in an adhoc fashion
- Need to support and abstract more use cases for optimization and transformation

# Planned Analyses and Transformations Framework

---

- DSL author defines transformation rules as pattern match on IR nodes
  - As we have done when IR nodes are created
- Delite provides IR traversal patterns
  - bottom-up
  - top-down
  - ...



# DSL Tooling and IDEs

---

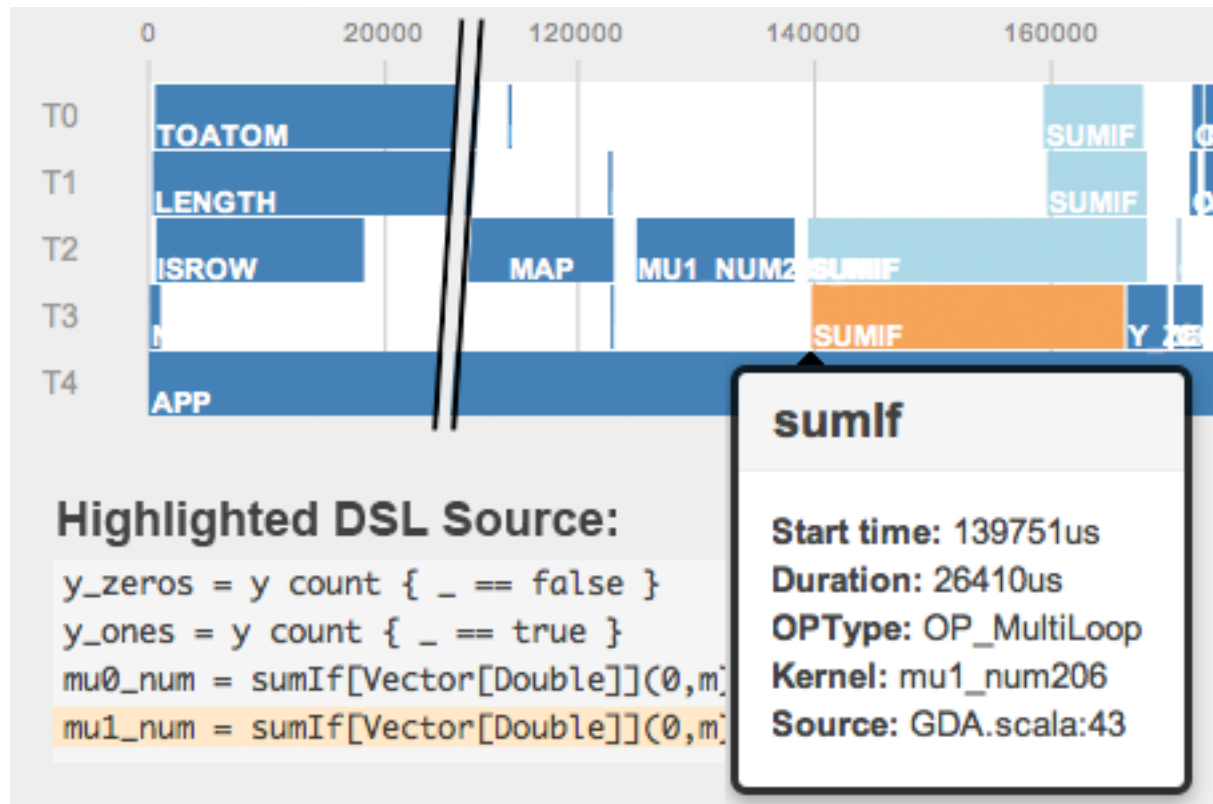
- With embedding can re-use host language but...
  - Very cryptic error messages
  - Refactoring support and other IDE goodies are missing
  - Difficult to debug generated code and trace it back to original DSL code

# Debugging Support for DSLs

---

- With DSLs, you can code at very high-level, it should be possible to debug at a very high level
  - Domain-specific debuggers
  - Domain-specific visualizer

# Debugging (2)



# Host Languages for DSLs

---

- Shown how we can modify Scala Compiler and make the language more overloadable and useful for high-performance embedding
- Still some issues
  - Can't overload class definition
  - Can't overload exceptions
- What would a host language built from the ground up for high-performance DSL embedding look like?